

Hashing

Abhiram Ranade

March 9, 2018

Concluding remarks on balanced search trees

Balanced search tree: Data structure to represent sets.

Operations allowed on a set S :

- ▶ $S.insert(x)$: Insert x into S .
- ▶ $S.find(x)$: Determine if x is present in S .
- ▶ $S.erase(x)$: Remove x from set.

All operations can be performed in time $O(\log n)$ where n is the number of elements currently in the set.

$O(\log n)$ time support also for $S.lower_bound(v)$ operation which returns **iterator** to element equal or larger than v .

Similarly upper bound operation.

Iterators can be used for many operations including scanning through elements in amortized $O(1)$ time per element.

Available in C++ STL class **set**

Maps

Map = function, set of pairs (x,y)

“(key,value)”
C++ class map

Operations allowed on a map M:

- ▶ $M.insert(x,y)$: insert the pair (x,y) into M.

Syntax: $M[x] = y;$

- ▶ $M.find(x)$: decide whether a pair (x,y) exists in M and if yes return y.

Syntax: $(M.count(x) > 0) ? M[x] : ...$

- ▶ $M.erase(x)$: delete pair (x,y) if it exists.

Other operations such as `upper_bound`, `lower_bound` also available.

Can scan through elements in increasing order of x using iterators.

Implemented using balanced search trees.

- ▶ Each node stores both x,y . Nodes are ordered by x .
- ▶ Insert, find, erase, upper/lower bound: $O(\log n)$ time.
- ▶ Scanning through elements in $O(1)$ amortized time.

Hashing: a faster way to represent sets and maps

- ▶ Insert, find, delete take time " $O(1)$ ".
Typically the time will be constant. Explained later.
- ▶ Lower and upper bound operations not supported.
Not needed in many applications.
- ▶ Available in C++ classes `unordered_set` and `unordered_map`.
key/set elements need not satisfy any ordering
Ordering relationship is not used in implementation.

A map (x,y) which supports `insert(x,y)`, `find(x)`, `erase(x)` is called a **dictionary**

Balanced search trees implement a dictionary such that each operation takes $O(\log n)$ time.

Hashing can be used to implement a dictionary such that each operation takes $O(1)$ time typically.

Warmup: Representing sets over a small universe

Suppose we wish to represent sets which are subsets of $\{0, \dots, 99\}$.

Implementation:

- ▶ For each set S maintain an array $s[0..99]$.
- ▶ Invariant: $s[i] == 1$ iff $i \in S$.
- ▶ $S.\text{insert}(i) : s[i] = 1$.
- ▶ $S.\text{find}[i] : s[i] == 1$
- ▶ $S.\text{erase}[i] : s[i] = 0$

“Direct address table”

Representing maps (x,y) where $x \in \{0, \dots, 99\}$

- ▶ Use an additional data array to store y part.
- ▶ $M.\text{insert}(x,y) : s[x] = 1; \text{data}[x] = y;$

What do we do if the key, x , is not from a small universe?

Example: 10 letter names: Table size = $26^{10} > 2^{45}$

Impractical!

Hashing: Approximating a direct address table

Suppose keys come from a universe $U = \{0, 1, \dots, |U| - 1\}$

Want to store $X \subset U$ with $|X| = n$.

Goal: Use $O(n)$ storage, and not $O(|U|)$.

Idea:

- ▶ Use a table T of size $m = kn$ for some small k .
- ▶ Select $h : U \rightarrow \{0, 1, \dots, m - 1\}$. e.g. $h(x) = x \bmod m$
- ▶ “If $x \in X$, $T[h(x)] = 1$ ”. Almost...
- ▶ What if $h(x) = h(y)$ for two keys $x \neq y$?
- ▶ $T[q] = \text{vector}/(\text{pointer to}) \text{ list of keys } x \text{ s.t. } h(x) = q$.
- ▶ Find(x) : Check if list (starting at) $T[h(x)]$ contains x .
- ▶ Delete(x) : Remove x from list (starting at) $T[h(x)]$

10 letter names: $U = \{10 \text{ digit number in radix } 26\}$

$x = x_0 + R(x_1 + R(x_2 + R(\dots)))$, where $R = 26$

To evaluate $h(x)$ perform each addition/mult mod m

Total storage used (table + names) $\approx kn + n \log |U|$.

Performance of hashing - “Typical case”

- ▶ Keys come from universe $U = \{0, \dots, |U| - 1\}$
- ▶ Table has size $m = kn$.
- ▶ Key mapping function $h(x) = x \bmod m$
- ▶ Some set X of n keys stored in table.

Say $k = 2$.

Assumption: We are storing n randomly drawn keys into X

- ▶ Keys will distribute nicely among table entries.
- ▶ $E[\text{number of keys in } T[i]] = O(n/m) = O(1)$
- ▶ Expected insertion time : $O(n/m) = O(1)$
- ▶ Expected find/delete time : $O(n/m) = O(1)$

$k = \frac{n}{m}$: “load factor” of table.

What if we don't know n ?

We resize, like vectors: Exercise

Performance of hashing - Worst case

- ▶ Keys come from universe $U = \{0, \dots, |U| - 1\}$
- ▶ Table has size m .
- ▶ Key mapping function $h(x) = x \bmod m$
- ▶ Some set X of n keys stored in table.

Worst case:

All $x \in X$ map to same table slot, i.e. $h(x) = \alpha$ for all $x \in X$.

Keys x, y have $h(x) = h(y)$: “collision”.

Each insertion will insert into the same list.

We should first check if the set already contains the key

$\Rightarrow n$ Insertions will take time $O(n^2)$

If we wish to *find*(y) where $h(y) = \alpha$, time = $O(n)$.

Is this likely?

Remarks

- ▶ Key requirement: Function h should distribute keys of interest over the table slots. “ h should appear to distribute randomly”
hash = random mess, so hash function h , hash table T
- ▶ Reasonably simple choices work for h . $h \bmod m$ will not work if keys are likely to be separated by αm .

Choose $m = \text{random prime!}$

$m = 2^k$, $h(x) = x \cdot A \bmod m$, where A is odd.

Something like this works for arbitrary sets w.h.p.

- ▶ Chaining: Handling collisions by keeping a list.
- ▶ Linear probing: If collision at $h(x)$ check if $h(x) + 1$ is empty, if not then $h(x) + 2$ and so on until an empty slot is found.
- ▶ Many other ways of resolving collisions also studied.
- ▶ Hash functions are typically much better than balanced trees if you only want insert/find/delete, and not lower/upper bound operations.
- ▶ C++ classes `unordered_set` and `unordered_map` provide iterators which can be used to step through all elements; but no order is guaranteed.

Other uses of hash functions

Minimizing transmission cost:

- ▶ Suppose you have file f_1 , your friend has file f_2 .
- ▶ You wish to know if $f_1 = f_2$.
- ▶ Can this be checked without transmitting entire files, if small probability of error is allowed?

Solution using agreed upon hash function h

- ▶ Your friend computes $h(f_2)$ and sends to you.
- ▶ You check if received $h(f_2) = h(f_1)$.
- ▶ If files are different received value will be different with large probability. Repeat for several hash functions.
Even single character difference will be detected.
- ▶ Transmission cost is small because hash value is small.

Error correction: When transmitting file f , send f and also “checksum” $c = h(f)$

Receiver checks that received f', c' satisfy $h(f') = c'$.